

# File System of Linux

Chonnam National University  
School of Electronics and Computer  
Engineering

Kyungbaek Kim

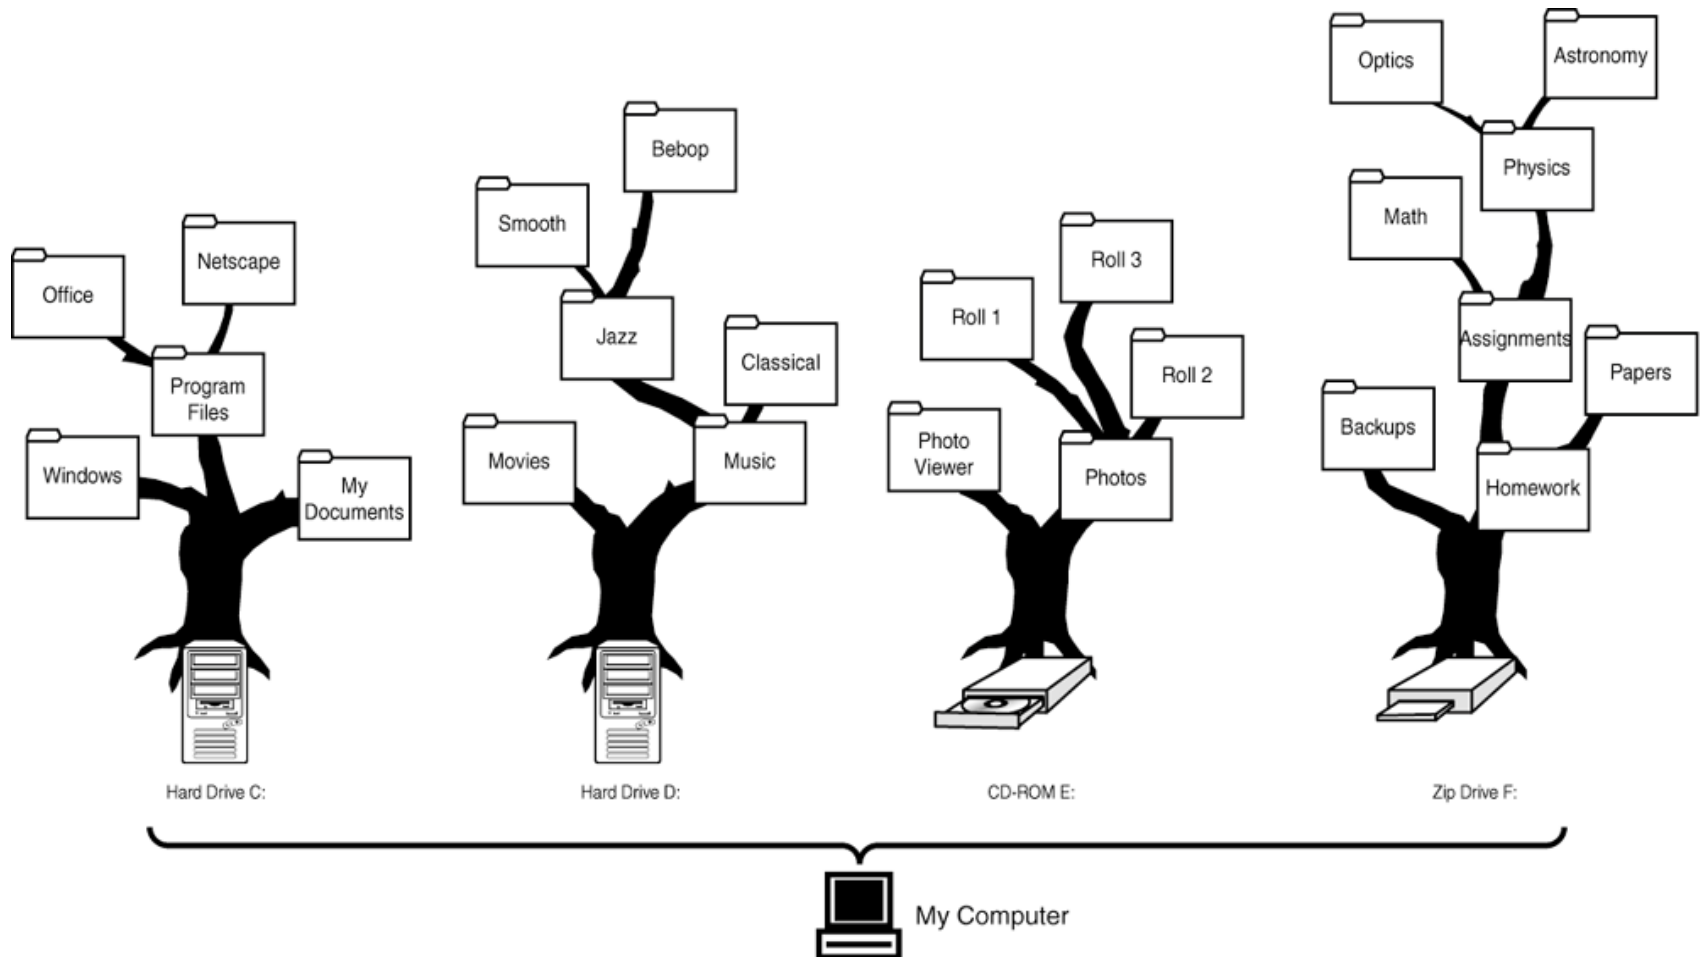
# Filesystems

- Commonly used to refer to two distinct concepts
  - Unified filesystem
    - The hierarchy of directories and files which humans use to organize data on a system
  - Filesystem types
    - The formatting system which the kernel uses to store blocks of data on physical media such as disks

# Unified Filesystem

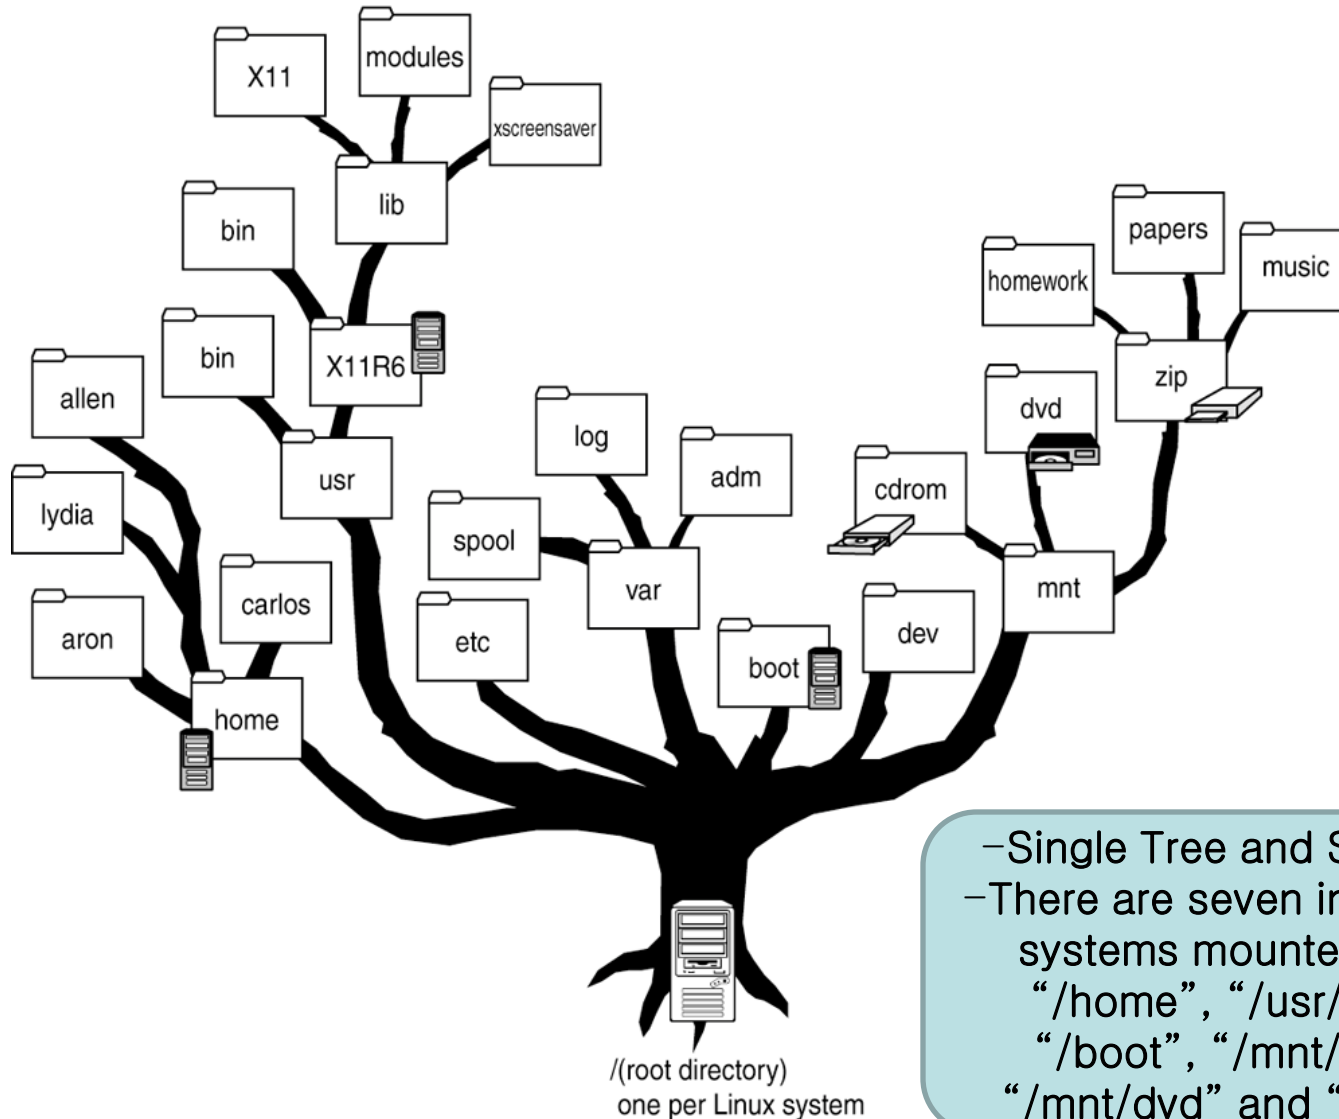
- Linux systems have a unified filesystem
  - Any file, on any disk or network share, can be accessed through a name beginning with “/”
    - The root
- Unified filesystem is made up of one or more **individual filesystems**
  - Each individual file system has its own root
  - That root can be grafted onto any directory in the unified filesystem
  - The directory where an individual filesystem is grafted into the unified filesystem is the individual filesystem's **mount point**
- An individual file system lives on a physical device(such as a disk drive), through not necessarily on the same computer

# MS Windows style File System



- Each device has its file system
- Many Trees and Many roots

# Linux Unified File System

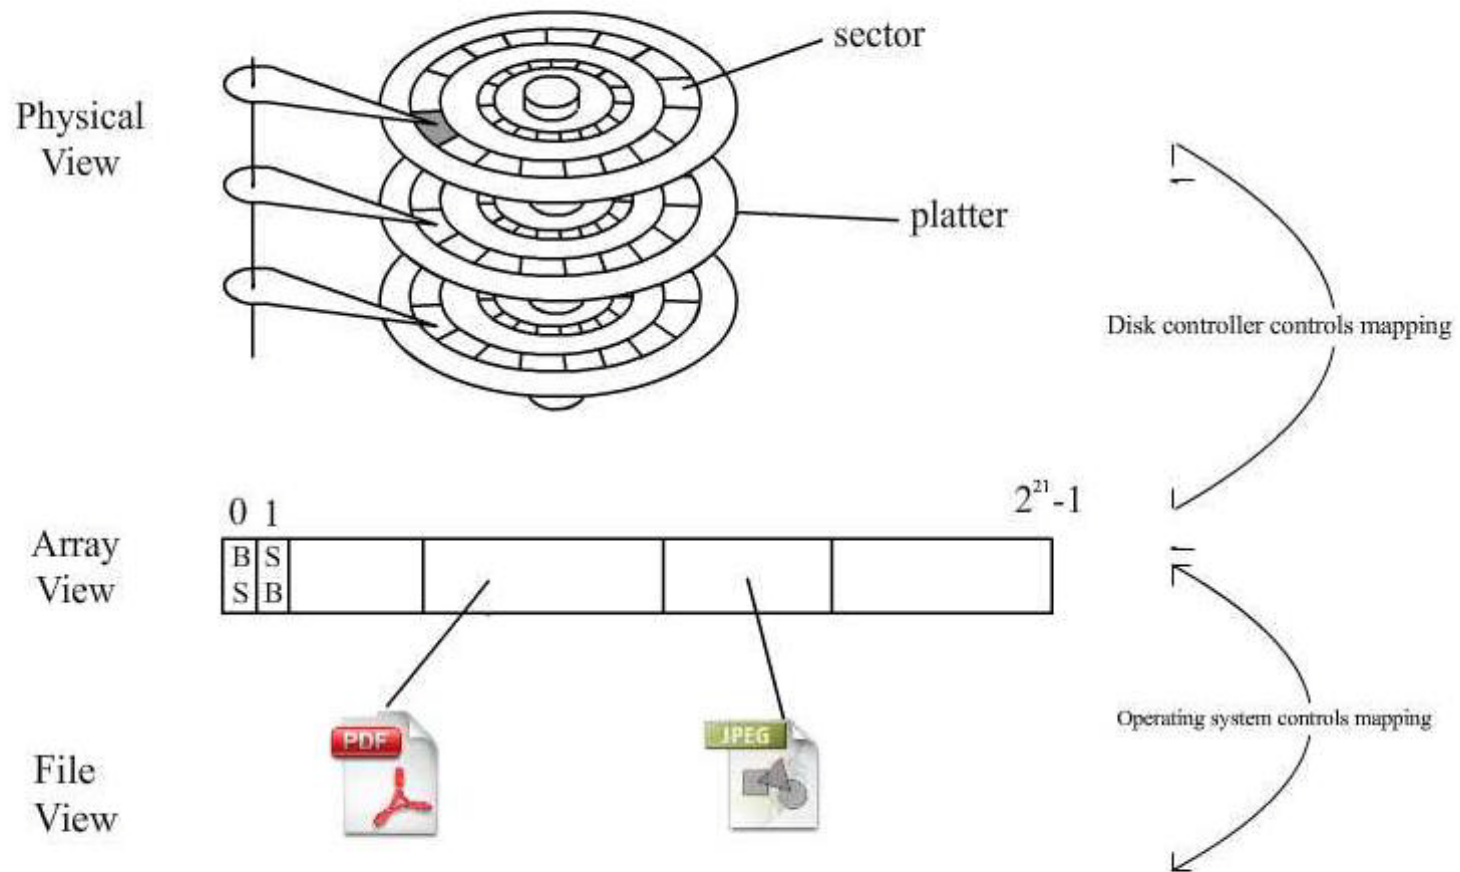


- Single Tree and Single Root
- There are seven individual file systems mounted on `/`, `/home`, `/usr/X11R6`, `/boot`, `/mnt/cdrom`, `/mnt/dvd` and `/mnt/zip`

# File Types

- Files directly contain data
- Directories provide a hierarchy of files: they contain both files and other directories
- *Files* and *directories* are both **file types**
- Other file types exist, including ***device special files***
  - Device files provide a way of asking the kernel for access to a given physical device
  - The data that the device file seems to contain is actually the raw sequence of bytes or sectors on the device itself
  - Device files are, by convention, stored under the **/dev** directory

# How to store data into your machine?



# Disk Naming

- A hard disk provides a single large storage space
- The device files for IDE hard drives are **/dev/hda**, **/dev/hdb**, **/dev/hdc** and **/dev/hdd**
  - hda and hdb are the drives on the first IDE channel, hdc and hdd the ones on the second channel
  - The first drive on each channel is the IDE master, and the second is the IDE slave
- SCSI (Small Computer System Interface) disks are named **/dev/sda**, **/dev/sdb**, etc

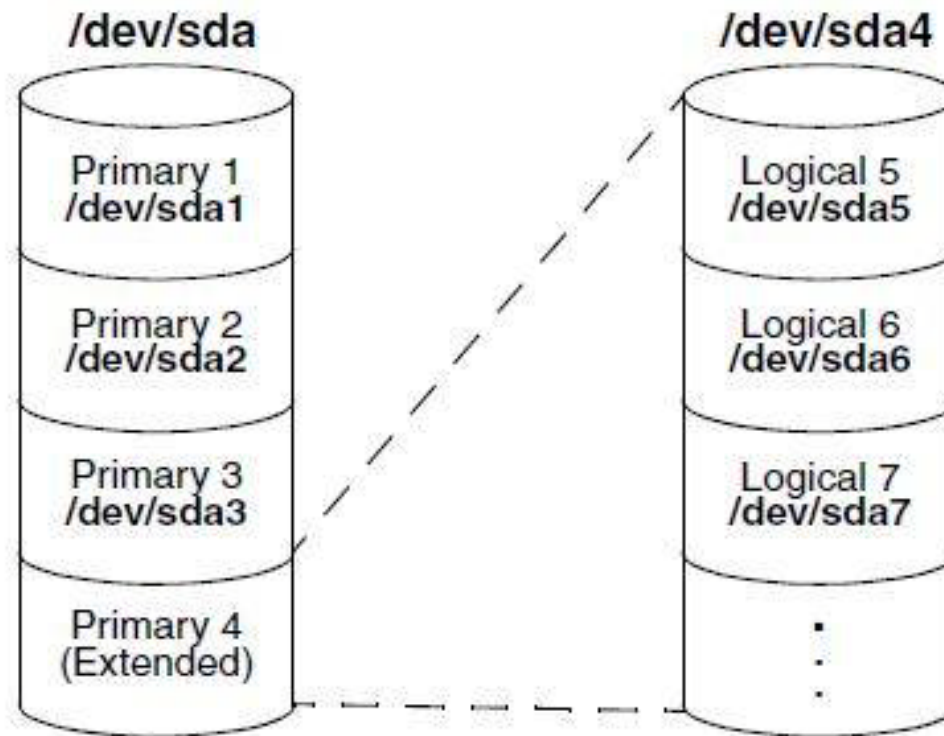




# Disks and Partitions

- Usually split into **partitions**
  - Information about partitions is stored in the **partition table**
  - Linux defaults using partition tables compatible with Microsoft Windows
  - For compatibility with Windows, at **most four primary partitions can be made**
    - *Primary partition → only one filesystem*
    - Primary partitions are numbered from 1–4
  - But they can be **extended partitions**, which can themselves be split into smaller **logical partitions**
    - Extended partitions have their own partition table to store information about logical partitions
    - Logical partitions are numbered from 5

# Primary vs Extended Partition



- The devices `/dev/hda`, etc., refer to whole hard disks, not partitions
  - Add the partition number to refer to a specific partition
  - `/dev/hda1` is the first partition on the first IDE disk

# Example of SCSI disks

```
kbkim@ubuntu: /dev
File Edit View Terminal Help

kbkim@ubuntu:/dev$ ls -l sda*
brw-rw---- 1 root disk 8, 0 2012-04-12 01:34 sda
brw-rw---- 1 root disk 8, 1 2012-04-12 01:34 sda1
brw-rw---- 1 root disk 8, 2 2012-04-12 01:34 sda2
brw-rw---- 1 root disk 8, 5 2012-04-12 01:34 sda5
kbkim@ubuntu:/dev$ ls -l disk/by-path/
total 0
lrwxrwxrwx 1 root root 9 2012-04-12 01:34 pci-0000:00:07.1-scsi-1:0:0:0 -> ../../sr0
lrwxrwxrwx 1 root root 9 2012-04-12 01:34 pci-0000:00:10.0-scsi-0:0:0:0 -> ../../sda
lrwxrwxrwx 1 root root 10 2012-04-12 01:34 pci-0000:00:10.0-scsi-0:0:0:0-part1 -> ../../sda1
lrwxrwxrwx 1 root root 10 2012-04-12 01:34 pci-0000:00:10.0-scsi-0:0:0:0-part2 -> ../../sda2
lrwxrwxrwx 1 root root 10 2012-04-12 01:34 pci-0000:00:10.0-scsi-0:0:0:0-part5 -> ../../sda5
kbkim@ubuntu:/dev$ ls -l disk/by-
by-path/ by-uuid/
kbkim@ubuntu:/dev$ ls -l disk/by-uuid/
total 0
lrwxrwxrwx 1 root root 10 2012-04-12 01:34 1905d52a-7b08-457a-86b0-ba472b882448 -> ../../sda5
lrwxrwxrwx 1 root root 10 2012-04-12 01:34 22a8bf2e-f825-45df-905b-cdb6c80231b3 -> ../../sda1
kbkim@ubuntu:/dev$
```

# fdisk command

- “fdisk” is used to create, delete and change the partitions on a disk
- “fdisk” reads one-letter commands from the user
  - Type “m” to get a list of commands
  - Use “p” to show what partitions currently exist
  - Use “q” to quit without alerting anything
  - Use “w” to quit and write the changes

# Command option of fdisk

```
root@ubuntu: /dev
File Edit View Terminal Help
root@ubuntu:/dev# fdisk sda

WARNING: DOS-compatible mode is deprecated. It's strongly recommended to
switch off the mode (command 'c') and change display units to
sectors (command 'u').

Command (m for help): m
Command action
  a toggle a bootable flag
  b edit bsd disklabel
  c toggle the dos compatibility flag
  d delete a partition
  l list known partition types
  m print this menu
  n add a new partition
  o create a new empty DOS partition table
  p print the partition table
  q quit without saving changes
  s create a new empty Sun disklabel
  t change a partition's system id
  u change display/entry units
  v verify the partition table
  w write table to disk and exit
  x extra functionality (experts only)

Command (m for help):
```

# Known Filesystem partition types

```
root@ubuntu: /dev
File Edit View Terminal Help
Command (m for help): l

0 Empty          24 NEC DOS       81 Minix / old Lin bf Solaris
1 FAT12          39 Plan 9       82 Linux swap / So c1 DRDOS/sec (FAT-
2 XENIX root     3c PartitionMagic 83 Linux           c4 DRDOS/sec (FAT-
3 XENIX usr      40 Venix 80286   84 OS/2 hidden C:  c6 DRDOS/sec (FAT-
4 FAT16 <32M     41 PPC PReP Boot 85 Linux extended  c7 Syrinx
5 Extended       42 SFS          86 NTFS volume set da Non-FS data
6 FAT16          4d QNX4.x        87 NTFS volume set db CP/M / CTOS / .
7 HPFS/NTFS      4e QNX4.x 2nd part 88 Linux plaintext de Dell Utility
8 AIX            4f QNX4.x 3rd part 8e Linux LVM        df BootIt
9 AIX bootable   50 OnTrack DM     93 Amoeba          e1 DOS access
a OS/2 Boot Manag 51 OnTrack DM6 Aux 94 Amoeba BBT       e3 DOS R/O
b W95 FAT32       52 CP/M         9f BSD/OS          e4 SpeedStor
c W95 FAT32 (LBA) 53 OnTrack DM6 Aux a0 IBM Thinkpad hi eb BeOS fs
e W95 FAT16 (LBA) 54 OnTrackDM6     a5 FreeBSD        ee GPT
f W95 Ext'd (LBA) 55 EZ-Drive      a6 OpenBSD        ef EFI (FAT-12/16/
10 OPUS          56 Golden Bow    a7 NeXTSTEP       f0 Linux/PA-RISC b
11 Hidden FAT12   5c Priam Edisk    a8 Darwin UFS     f1 SpeedStor
12 Compaq diagnost 61 SpeedStor     a9 NetBSD         f4 SpeedStor
14 Hidden FAT16 <3 63 GNU HURD or Sys ab Darwin boot    f2 DOS secondary
16 Hidden FAT16   64 Novell Netware af HFS / HFS+     fb VMware VMFS
17 Hidden HPFS/NTF 65 Novell Netware b7 BSDI fs        fc VMware VMKCORE
18 AST SmartSleep 70 DiskSecure Mult b8 BSDI swap      fd Linux raid auto
1b Hidden W95 FAT3 75 PC/IX         bb Boot Wizard hid fe LANstep
1c Hidden W95 FAT3 80 Old Minix      be Solaris boot   ff BBT
1e Hidden W95 FAT1
```

# Partition Table

```
root@ubuntu: /dev
File Edit View Terminal Help

root@ubuntu:/dev# fdisk sda

WARNING: DOS-compatible mode is deprecated. It's strongly recommended to
switch off the mode (command 'c') and change display units to
sectors (command 'u').


Command (m for help): p

Disk sda: 21.5 GB, 21474836480 bytes
255 heads, 63 sectors/track, 2610 cylinders
Units = cylinders of 16065 * 512 = 8225280 bytes
Sector size (logical/physical): 512 bytes / 512 bytes
I/O size (minimum/optimal): 512 bytes / 512 bytes
Disk identifier: 0x000a1bb6

Device Boot      Start         End      Blocks   Id  System
sda1   *           1         2497     20051968   83   Linux
sda2             2497         2611       916481    5   Extended
sda5             2497         2611       916480   82   Linux swap / Solaris

Command (m for help):
```

# Sector/Track(head)/Cylinder





# mkfs command

- “mkfs” initializes a filesystem on a new partition
  - **Warning!! : any old data on the partition will be lost**
  - e.g.) `$ mkfs -t ext4 -c /dev/sda2`
    - Make an ext4 filesystem on /dev/sda2

# Types of filesystems

- FAT(File Allocation Table) : FAT12/16/32  
→ DOS, Windows 9x
- NTFS (Net Technology File System)  
→ Windows NT, windows 2000, windows XP/7
- HFS (Hierarchical File System), HFS+ → Mac
- HPFS (High Performance File system) → OS/2 operation system
- UFS (Unix File System) → BSD unix file system
- **ext2, ext3, ext4 (extended filesystem)**  
→ **Linux**, Android 2.3
- VMFS (Virtual Machine File System) → VMWare
- UDF (Universal Disk Format) → Optical disks

# Mounting Filesystems

- A partition contains entirely arbitrary data
- One filesystem is made **the root filesystem**
  - the root directory on that filesystem becomes the directory named “/”
- Other filesystems can **be mounted**
  - The root directory of that filesystem is grafted onto a directory of the root filesystems
  - This arranges for every file in every mounted filesystem to be accessible from a single unified name space
- The directory grafted onto is called the **mount point**

# mount

- Important filesystems are mounted at boot-up
  - Other filesystems can be mounted or unmounted at any time
- “mount” mounts a filesystem
  - You usually need to have root permission to mount a filesystem
- “mount” makes it easy to mount filesystems configured by the system administrator


# Mounting other filesystems

- “mount /dev/sdb3 /mnt/extra”
  - mounts the filesystem stored in the “/dev/sdb3” device on the mount point “/mnt/extra”
- You may occasionally need to specify the filesystem type explicitly
  - e.g.) \$ mount -t vfat /dev/hdd1 /mnt/windows

# umount

- “umount” unmounts a filesystem
  - Note the spelling.
- “umount /mnt/extra”
  - Unmounts whatever is on the /mnt/extra mount point
- “umount /dev/sdb3”
  - Unmounts any filesystem in the /dev/sdb3 device
- Need root permission
- It is also impossible to unmount a “busy” filesystem
  - If a process has file on it open
  - If a process has a directory with it as its current directory

# Example of mount/umount : setting for VMWare

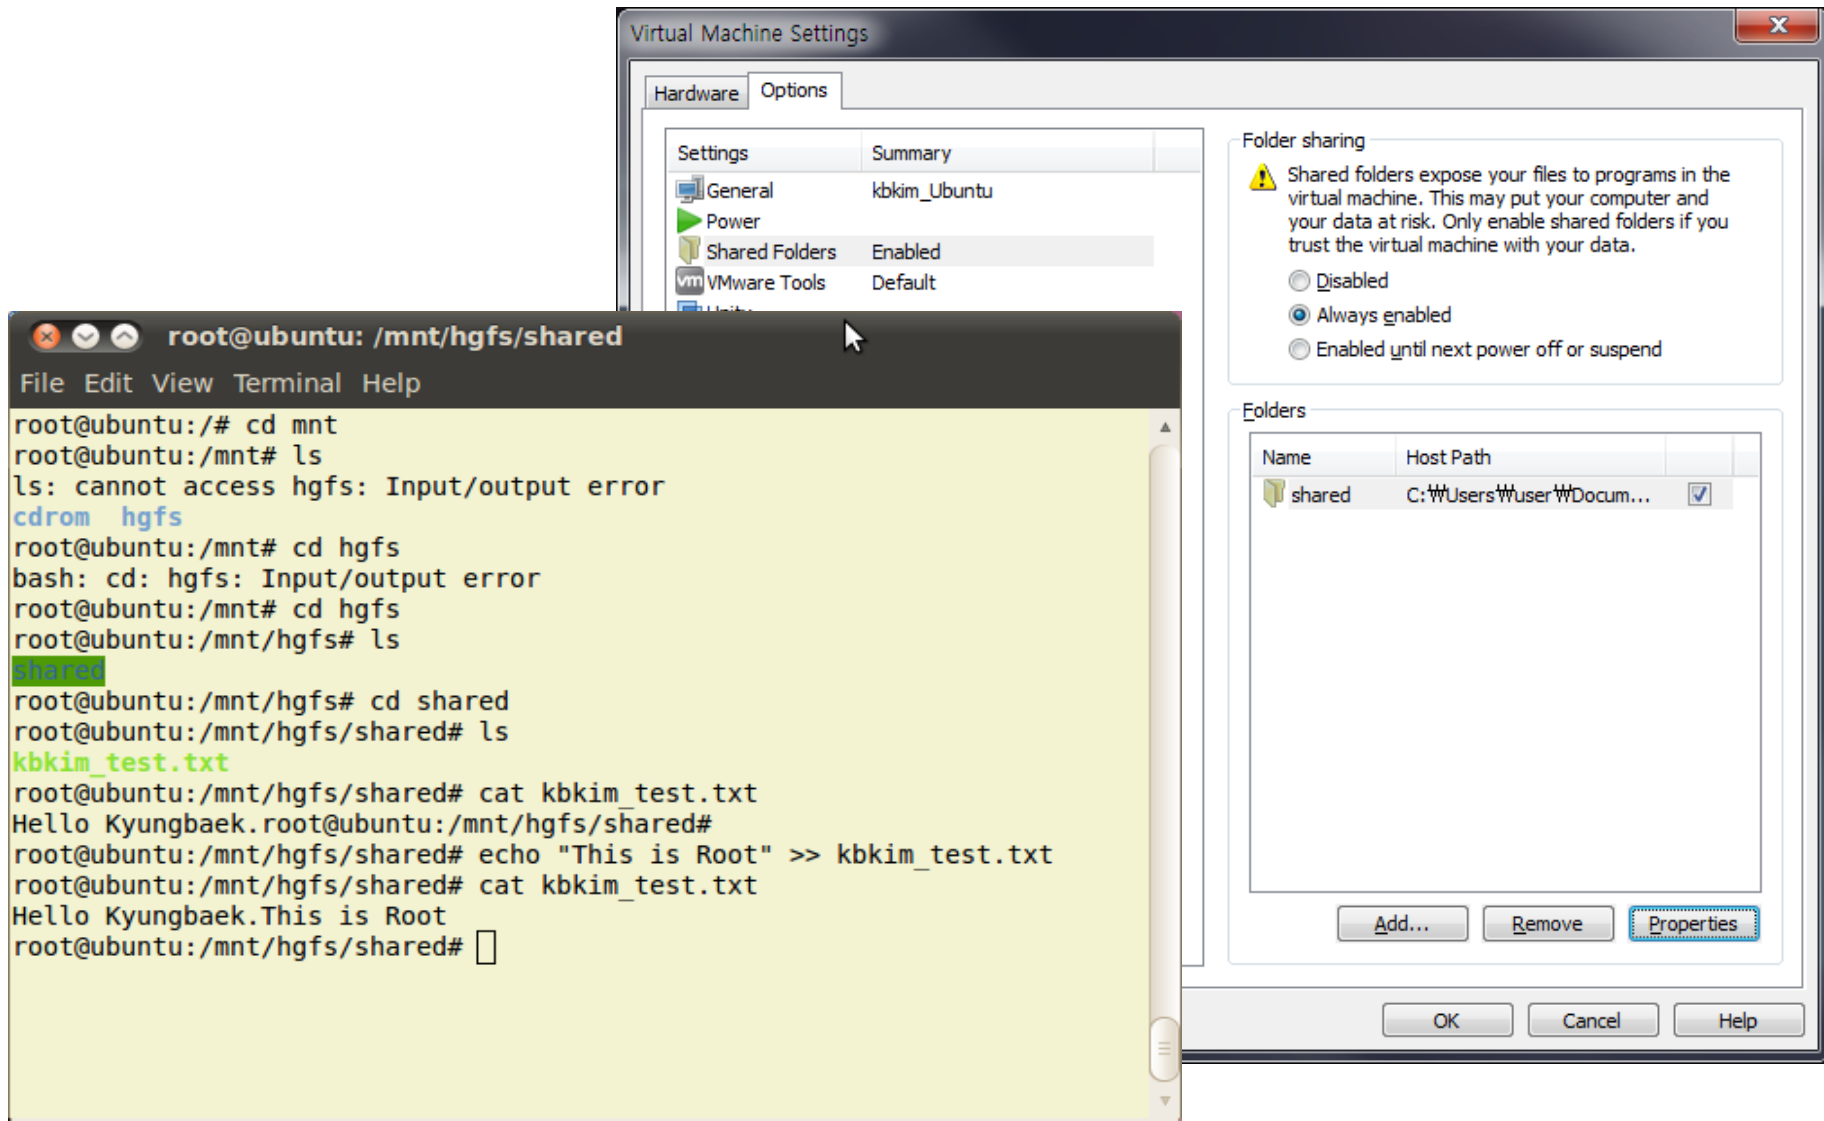


# Example of mount/umount

```
root@ubuntu: /mnt
File Edit View Terminal Help
root@ubuntu:/mnt# cd cdrom
root@ubuntu:/mnt/cdrom# ls
root@ubuntu:/mnt/cdrom# cd ..
root@ubuntu:/mnt# mount -r /dev/cdrom /mnt/cdrom
root@ubuntu:/mnt# cd cdrom
root@ubuntu:/mnt/cdrom# ls
autorun.inf  dists      isolinux   pics       preseed    ubuntu
casper      install   md5sum.txt pool       README.diskdefines  wubi.exe
root@ubuntu:/mnt/cdrom# umount /mnt/cdrom
umount: /mnt/cdrom: device is busy.
(In some cases useful info about processes that use
the device is found by lsof(8) or fuser(1))
root@ubuntu:/mnt/cdrom# cd ..
root@ubuntu:/mnt# umount /mnt/cdrom
root@ubuntu:/mnt# cd cdrom
bash: cd: cdrom: No such file or directory
root@ubuntu:/mnt#
```



# HGFS (Host Guest File System)



# Configuring mount : /etc/fstab

- “/etc/fstab” contains information about filesystems that are known the system administrator
  - Specifying a filesystem in /etc/fstab makes it possible to use its mount point as the only argument to mount
- Configure Boot-up mount
- Each line describes one filesystem
- Six columns on each line

# Example of fstab file

```
root@ubuntu: /mnt
File Edit View Terminal Help

root@ubuntu:/mnt# cat /etc/fstab
# /etc/fstab: static file system information.
#
# Use 'blkid -o value -s UUID' to print the universally unique identifier
# for a device; this may be used with UUID= as a more robust way to name
# devices that works even if disks are added and removed. See fstab(5).
#
# <file system> <mount point> <type> <options> <dump> <pass>
proc /proc proc nodev,noexec,nosuid 0 0
# / was on /dev/sda1 during installation
UUID=22a8bf2e-f825-45df-905b-cdb6c80231b3 / ext4 errors=remount-ro 0 1
# swap was on /dev/sda5 during installation
UUID=1905d52a-7b08-457a-86b0-ba472b882448 none swap sw 0 0
/dev/fd0 /media/floppy0 auto rw,user,noauto,exec,utf8 0 0
/dev/cdrom /mnt/cdrom auto ro 0 0
root@ubuntu:/mnt# mount /mnt/cdrom
root@ubuntu:/mnt# ls cdrom
autorun.inf dists isolinux pics preseed ubuntu
casper install md5sum.txt pool README.diskdefines wubi.exe
root@ubuntu:/mnt# umount /mnt/cdrom
root@ubuntu:/mnt#
```

# Mount options

- Comma-separated options in `/etc/fstab`
- Alternatively, use comma-separated options with `-o` on the mount command line
- Options
  - “auto” → mount a filesystem on boot-up
  - “noauto” → prevent mounting a filesystem on boot-up, useful for removable media
  - “ro” → read only
  - “rw” → read/write
  - “users” → let non-root users mount/umount

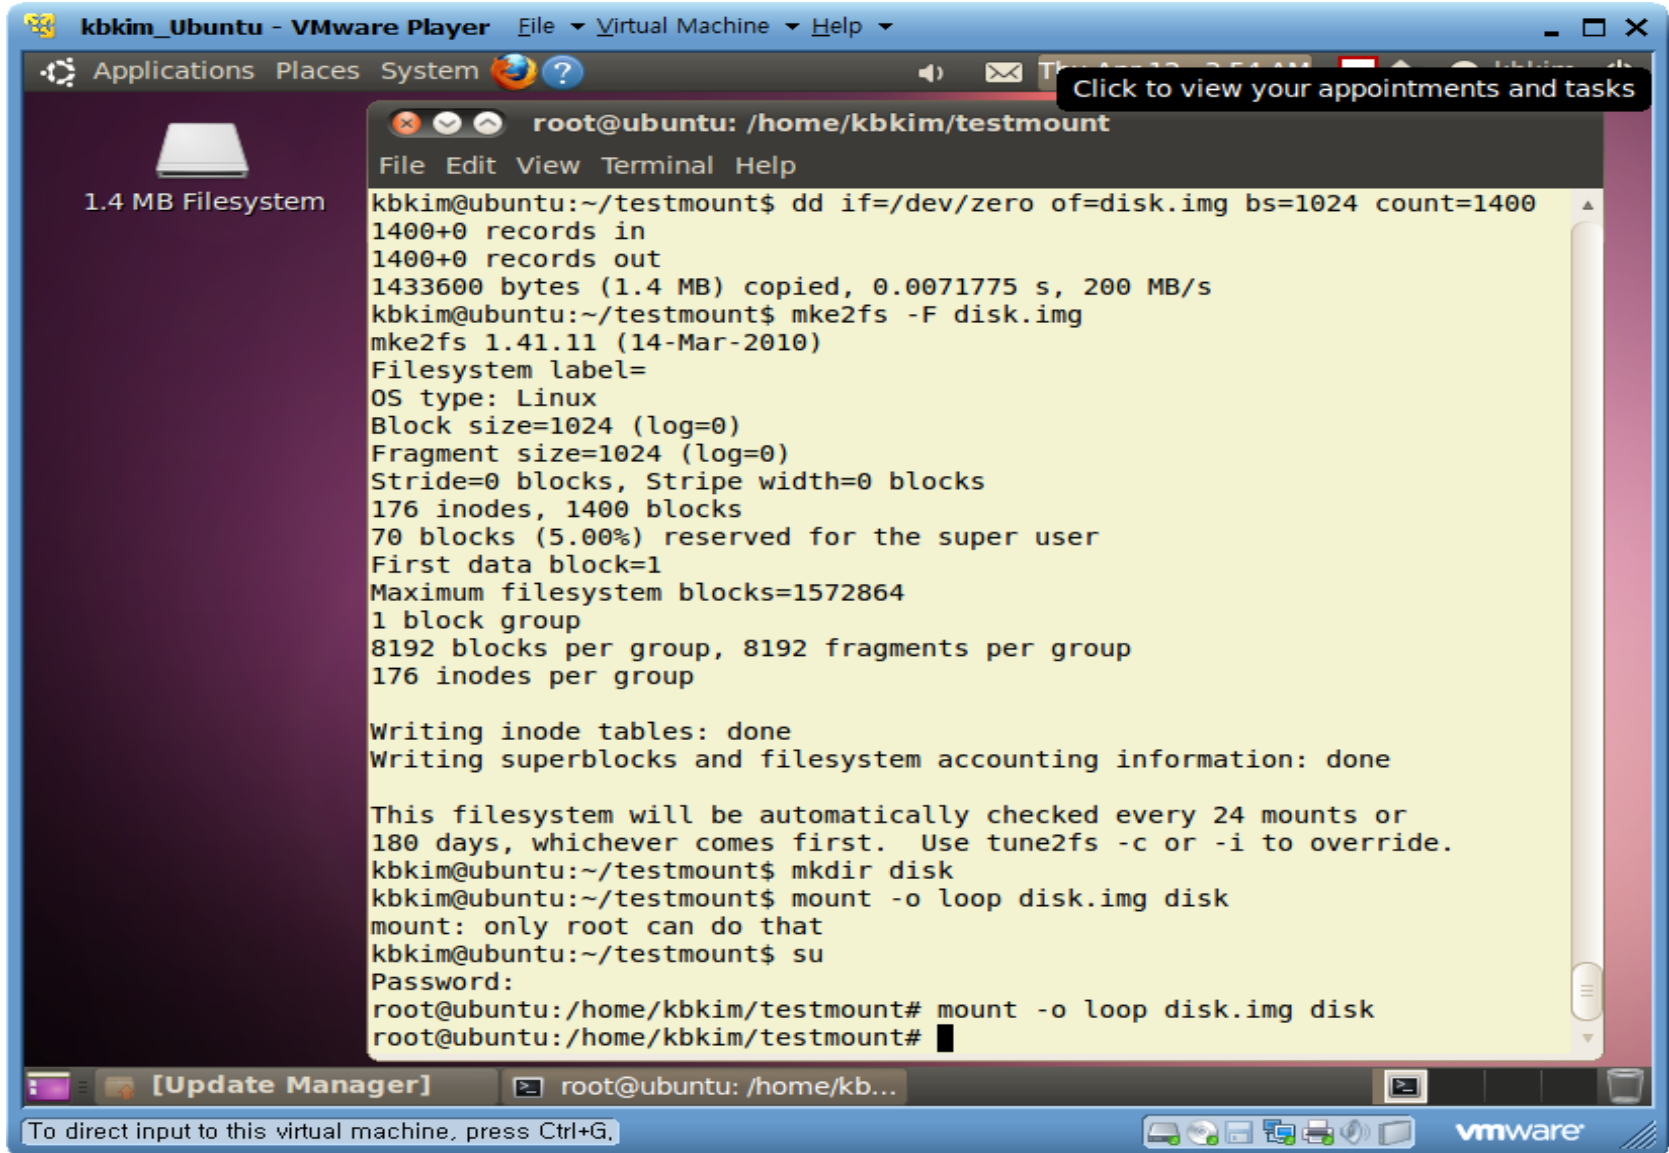
# Other columns in /etc/fstab

- Dump
  - Used by the dump and restore backup utilities
  - Just use 1 for normal filesystems and for 0 for removable filesystems
- Pass-no
  - Controls the order in which automatically-mounted filesystems are checked by fsck
  - Use 1 for the root filesystem
  - Use 0 for filesystems that are not mounted at boot-up
  - Use 2 for other filesystems

# Mounting a file

- Using loop devices
  - Linux can mount a filesystem stored in a normal file, instead of a disk
  - Use “-o loop” option
  - e.g.) image files
- Useful for testing images of CD-ROMs before burning them to disk

# Example of mounting a file



```
kbkim_Ubuntu - VMware Player  File  Virtual Machine  Help
Applications  Places  System  ?
1.4 MB Filesystem

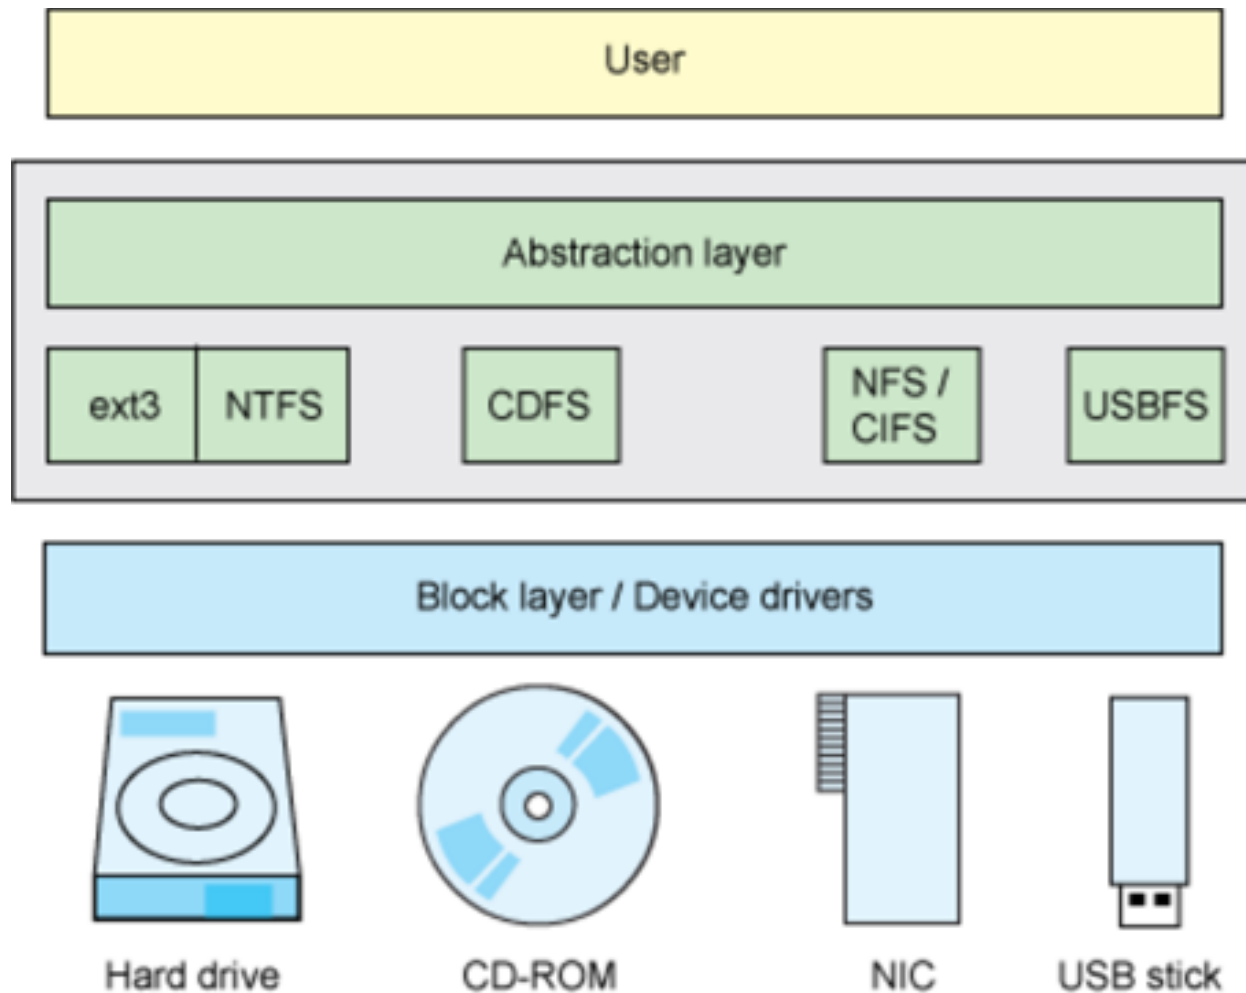
root@ubuntu: /home/kbkim/testmount
File Edit View Terminal Help
kbkim@ubuntu:~/testmount$ dd if=/dev/zero of=disk.img bs=1024 count=1400
1400+0 records in
1400+0 records out
1433600 bytes (1.4 MB) copied, 0.0071775 s, 200 MB/s
kbkim@ubuntu:~/testmount$ mke2fs -F disk.img
mke2fs 1.41.11 (14-Mar-2010)
Filesystem label=
OS type: Linux
Block size=1024 (log=0)
Fragment size=1024 (log=0)
Stride=0 blocks, Stripe width=0 blocks
176 inodes, 1400 blocks
70 blocks (5.00%) reserved for the super user
First data block=1
Maximum filesystem blocks=1572864
1 block group
8192 blocks per group, 8192 fragments per group
176 inodes per group

Writing inode tables: done
Writing superblocks and filesystem accounting information: done

This filesystem will be automatically checked every 24 mounts or
180 days, whichever comes first.  Use tune2fs -c or -i to override.
kbkim@ubuntu:~/testmount$ mkdir disk
kbkim@ubuntu:~/testmount$ mount -o loop disk.img disk
mount: only root can do that
kbkim@ubuntu:~/testmount$ su
Password:
root@ubuntu:/home/kbkim/testmount# mount -o loop disk.img disk
root@ubuntu:/home/kbkim/testmount#
```

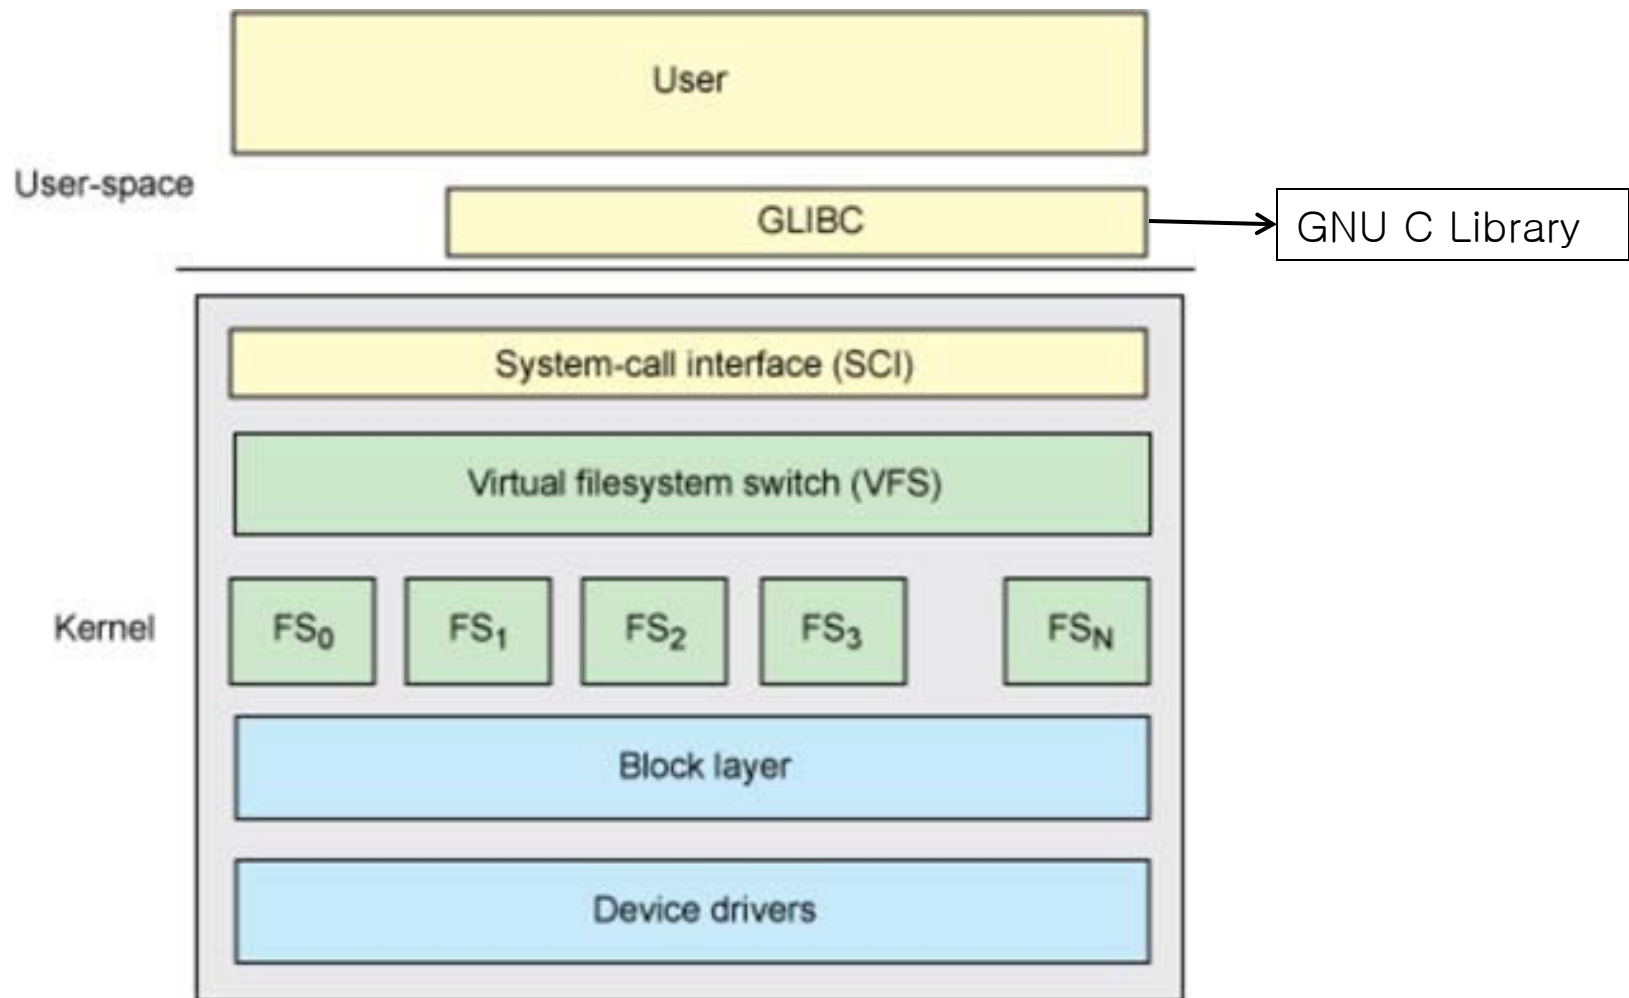
[Update Manager] root@ubuntu: /home/kb...  
To direct input to this virtual machine, press Ctrl+G.

# Abstraction of Multiple File System





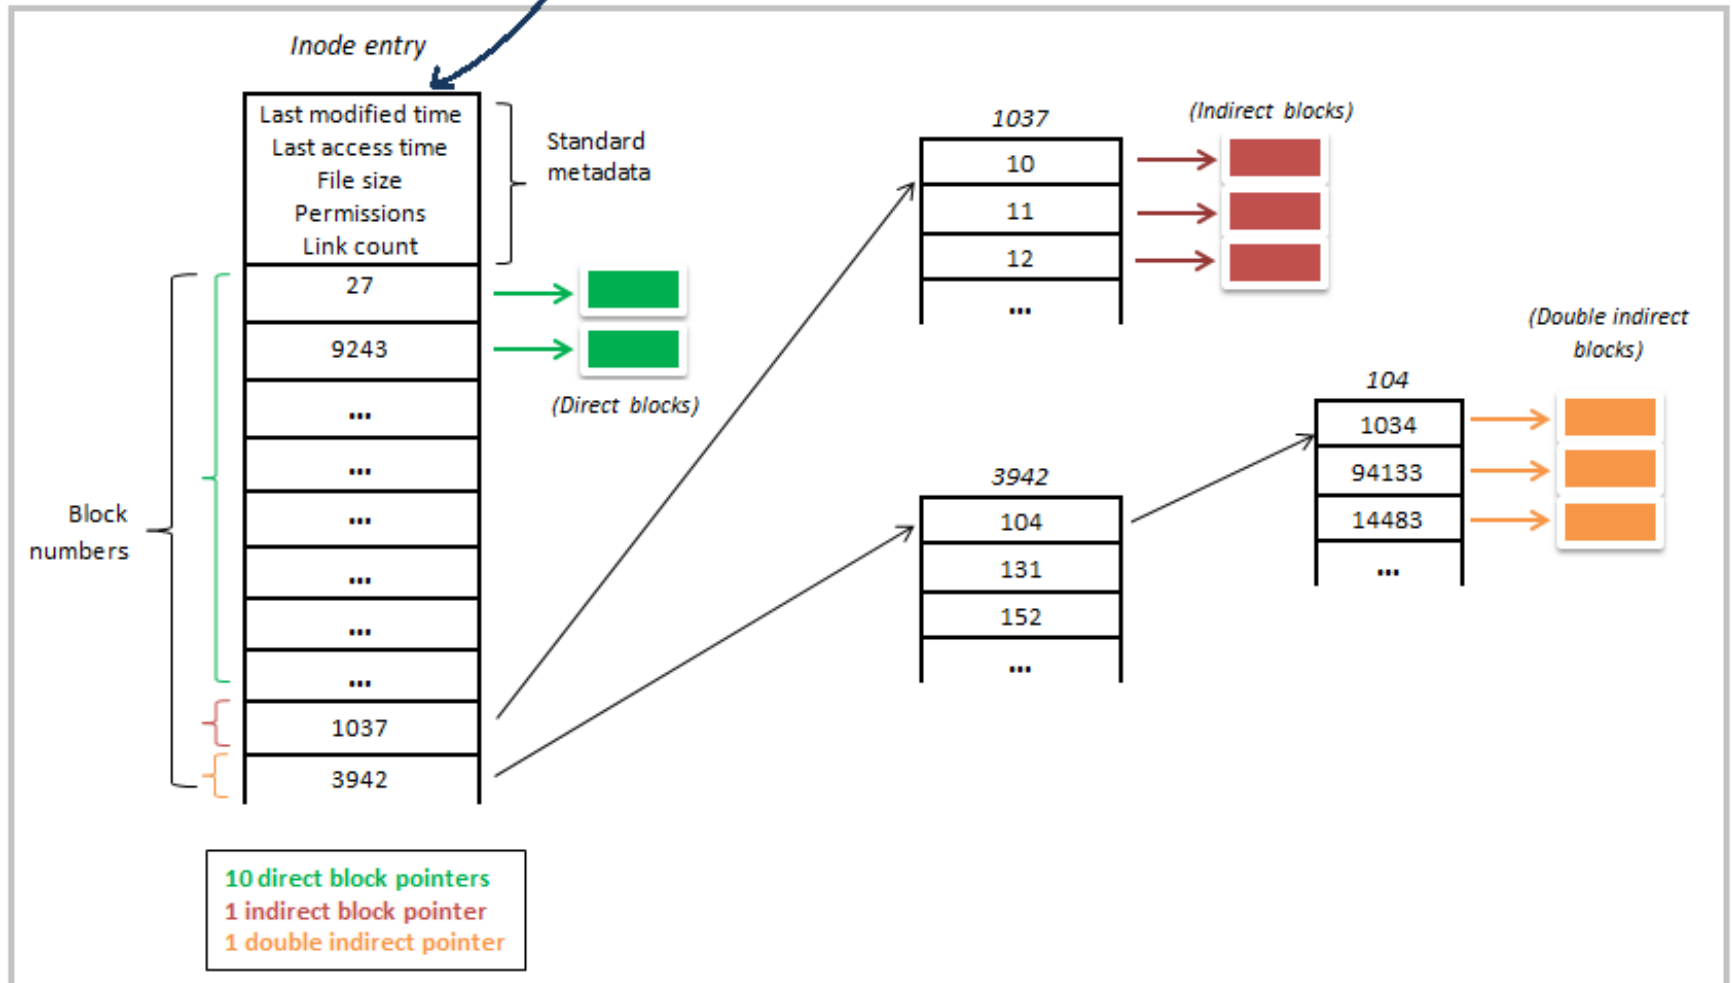

# Virtual Filesystem Switch (VFS)



# Inodes (Index nodes)

- Data structure that describes a file on an individual filesystem
- It contains information about the file, including its type (file/directory/device), size, modification, time, permissions, etc
- You can regard an inode as being the file itself
- The Inodes within an individual filesystem are numbered → inum

## Unix File System




# Directory

- File's name is stored not in its inode, but in directory
- Stored on disk as a list of file and directory names
- Each name has an inode number associated with it
- Separating names from Inodes
  - multiple directory entries referring to the same file

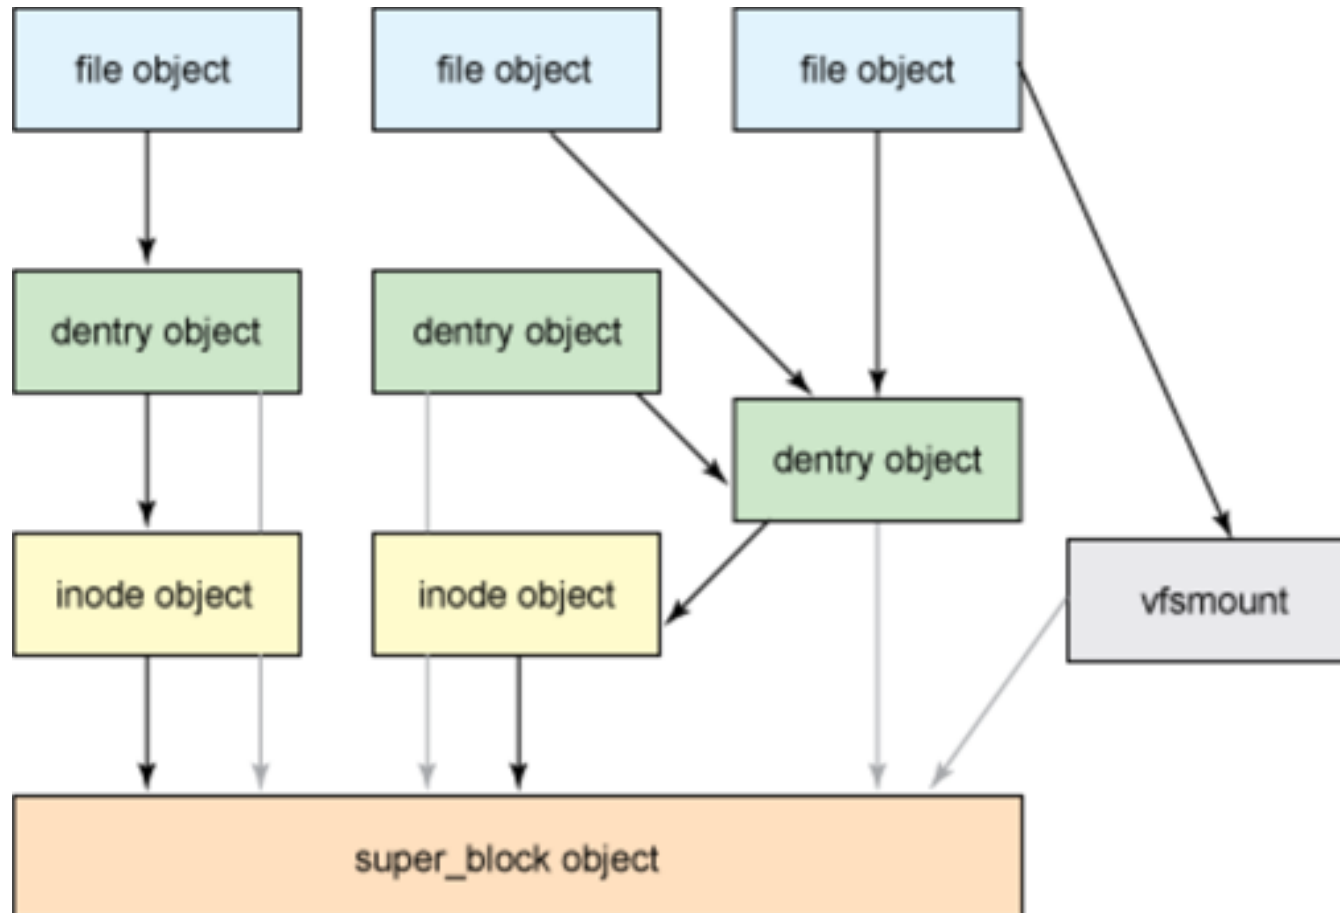
	inode	rec_len	file_type	name_len	name	
0	13	12	1	2	.	\0 \0 \0
12	2	12	2	2	..	\0 \0
24	18	16	5	2	m u s i c	\0 \0 \0
40	15	16	8	1	t e s t . t x t	
56	19	12	3	2	b i n	\0

## The structure of a directory entry



## Resolution of file path

# In-memory VFS objects



### struct file

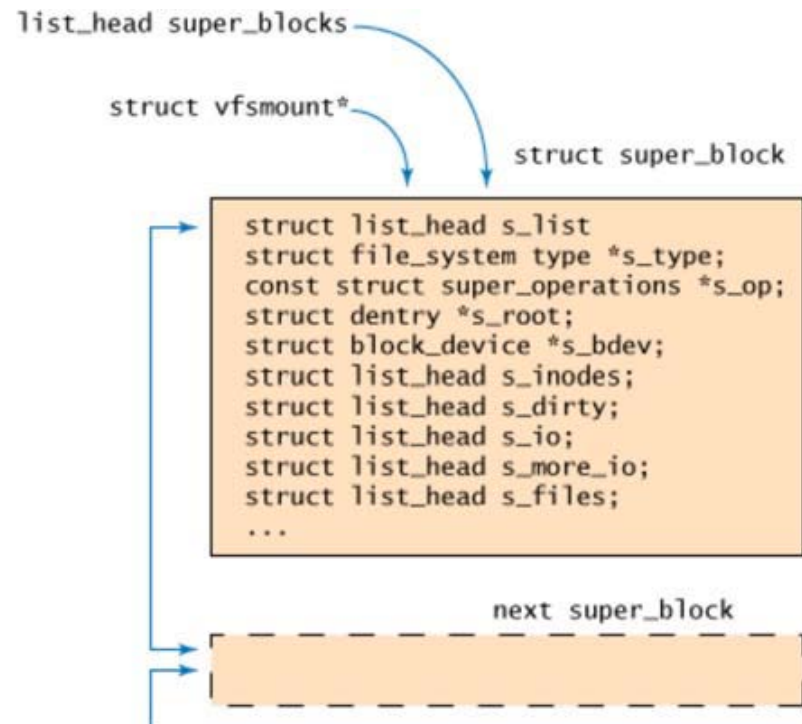
```
struct path f_path;  
struct dentry (f_path.dentry);  
const struct file_operations *f_op;  
unsigned int f_flags;  
fmode_t f_mode;  
loff_t f_pos;  
...
```

### struct dentry

```
struct super_block *d_sb;  
struct dentry *d_parent;  
struct list_head d_subdirs;  
struct dentry_operations *d_op;  
unsigned char d_iname[];  
struct inode *d_inode;  
...
```

### struct inode

```
struct list_head i_dentry;  
struct timespec i_atime;  
struct timespec i_mtime;  
struct timespec i_ctime;  
gid_t i_gid;  
uid_t i_uid;  
loff_t i_size;  
const struct file_operations *i_fop;  
const struct inode_operations *i_op;  
struct address_space *i_mapping;  
struct address_space *i_data;  
...
```



### Superblock

- Store the information of a filesystem
- Essential to mount a filesystem
- List of `super_block` means that multiple filesystems are in a device

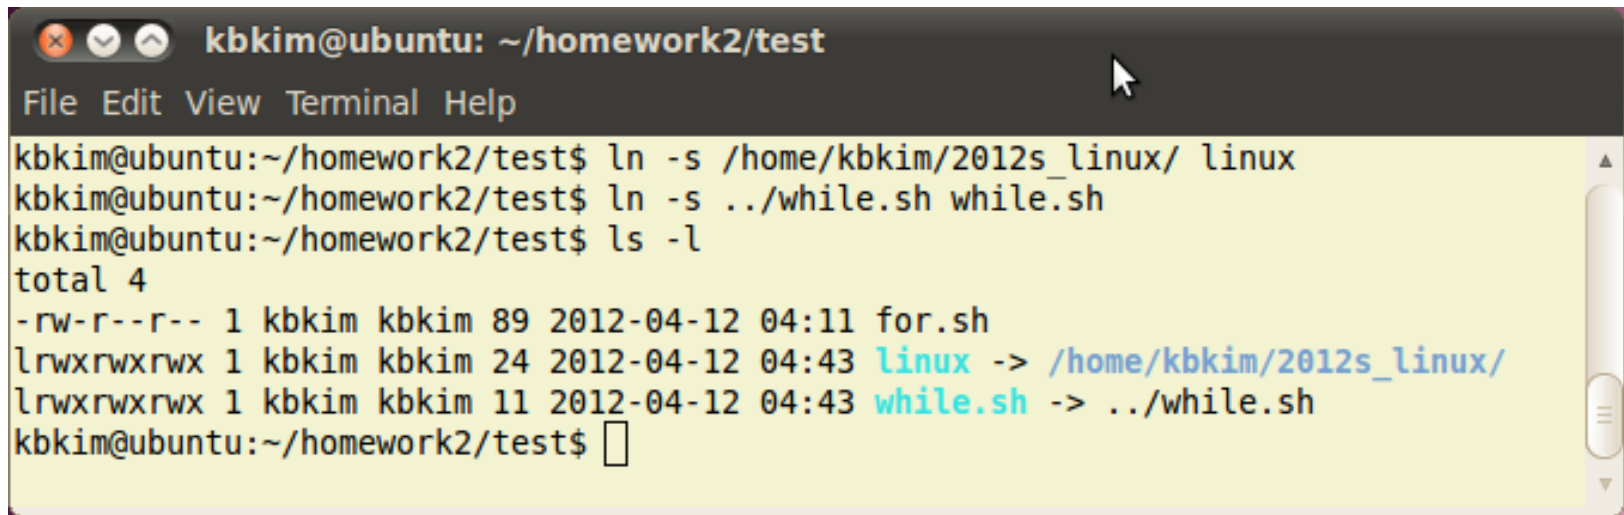
# Symbolic Links

- A symbolic link (or symlink) is a pseudo-file which behaves as an **alternative name** for some other file or directory
- The contents of the symlink are the real name pointed to
  - Kernel replace the symlink contents automatically
- **Keep a file in one place, but pretend it lives in another**
  - To ensure that an obsolete name continues to work for older software
  - To spread data from a single filesystem hierarchy over multiple disk partitions



# Examining and creating symbolic links

- “ls -l” shows where a symbolic link points to
  - Different color or suffix with ‘@’
- “ln -s” create a symbolic link



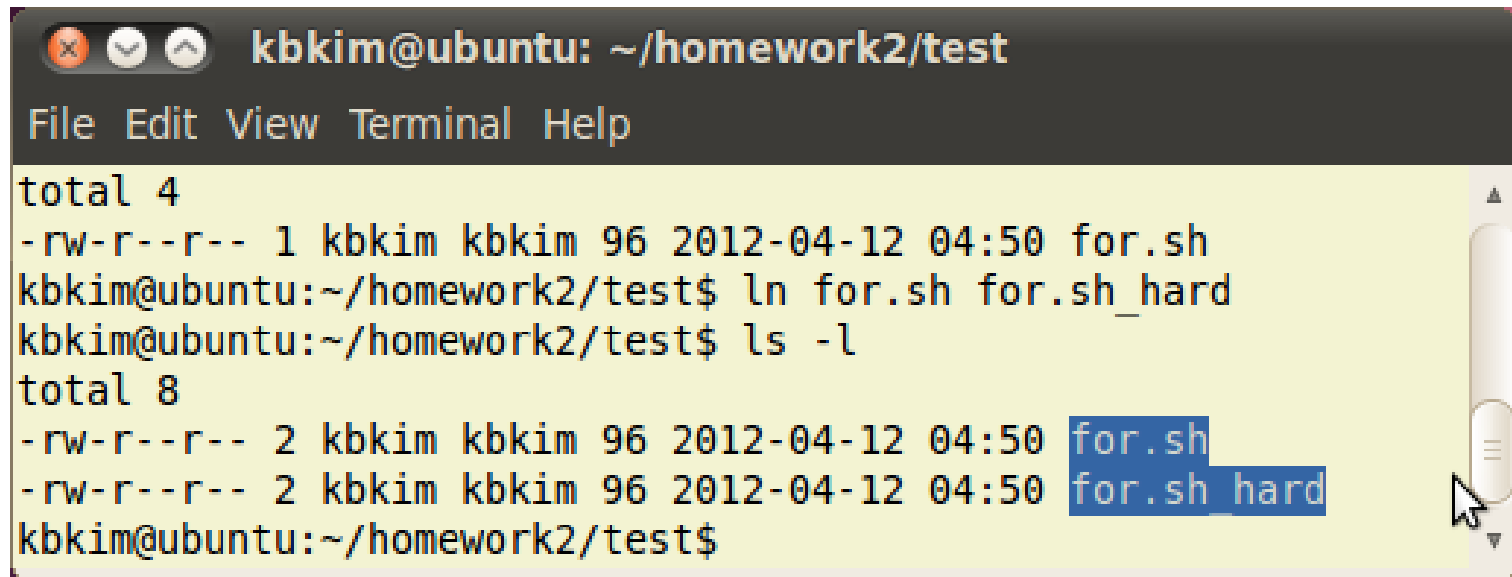
```
kbkim@ubuntu: ~/homework2/test
File Edit View Terminal Help
kbkim@ubuntu:~/homework2/test$ ln -s /home/kbkim/2012s_linux/ linux
kbkim@ubuntu:~/homework2/test$ ln -s ../while.sh while.sh
kbkim@ubuntu:~/homework2/test$ ls -l
total 4
-rw-r--r-- 1 kbkim kbkim 89 2012-04-12 04:11 for.sh
lrwxrwxrwx 1 kbkim kbkim 24 2012-04-12 04:43 linux -> /home/kbkim/2012s_linux/
lrwxrwxrwx 1 kbkim kbkim 11 2012-04-12 04:43 while.sh -> ../while.sh
kbkim@ubuntu:~/homework2/test$
```

# Hard Links

- Hard link refers to another file by inode number
- A directory entry contains a name and an inode number
- **The modification of file affected to both of the original file and the hard link file**
  - But deleting or renaming one does not affect the other

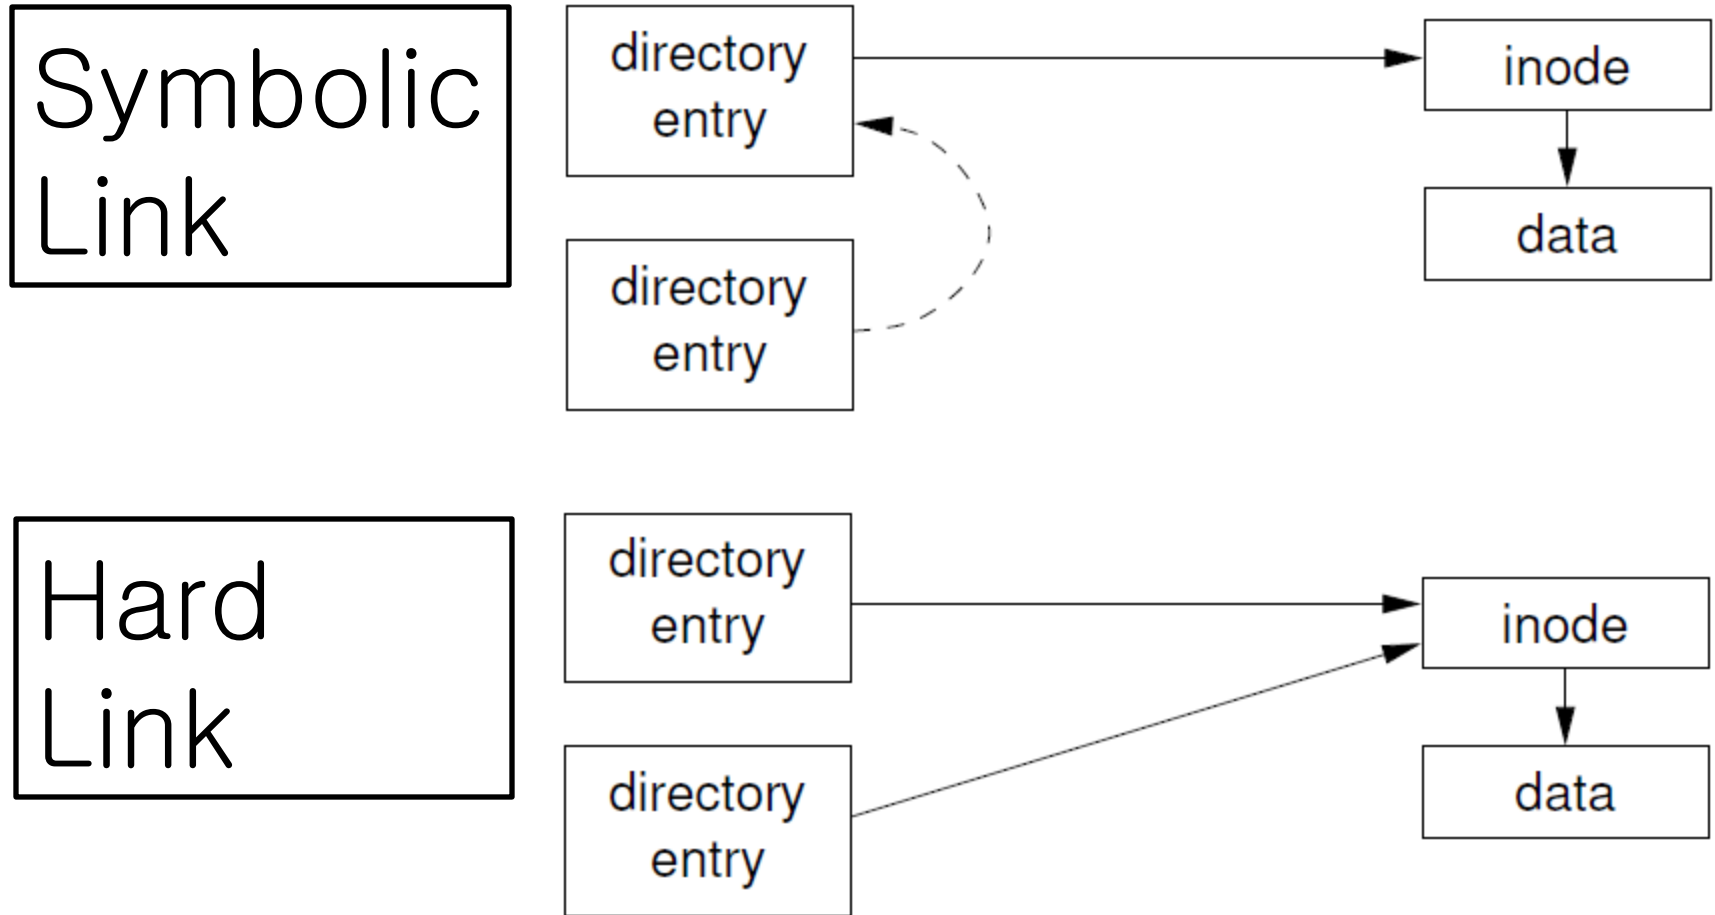
# Examining and creating hard links

- “ls” shows a hard link file has multiple link count
- “ln” create a hard link

A terminal window titled 'kbkim@ubuntu: ~/homework2/test' with a menu bar (File, Edit, View, Terminal, Help). The terminal shows the output of 'ls' (total 4, for.sh) followed by the command 'ln for.sh for.sh\_hard'. A second 'ls -l' command shows two files, 'for.sh' and 'for.sh\_hard', both with a link count of 2. The file names 'for.sh' and 'for.sh\_hard' are highlighted in blue in the original image.

```
kbkim@ubuntu: ~/homework2/test
File Edit View Terminal Help
total 4
-rw-r--r-- 1 kbkim kbkim 96 2012-04-12 04:50 for.sh
kbkim@ubuntu:~/homework2/test$ ln for.sh for.sh_hard
kbkim@ubuntu:~/homework2/test$ ls -l
total 8
-rw-r--r-- 2 kbkim kbkim 96 2012-04-12 04:50 for.sh
-rw-r--r-- 2 kbkim kbkim 96 2012-04-12 04:50 for.sh_hard
kbkim@ubuntu:~/homework2/test$
```

# Symbolic Link vs Hard Link



# Symbolic link vs Hard link

Symbolic links	Hard Links
Distinguishable (pointer)	indistinguishable
Point to any type of file	may not point to a directory
Refer to names	Works by inode numbers
Point on other filesystems	Work within a single filesystem
Renaming or deleting a original file breaks the symbolic link	Renaming or deleting a original file has no effect on the hard link
Take up additional disk space (to store the name pointed to)	Need as much disk space as a directory entry

# Inode and symbolic/hard link

```
kbkim@ubuntu: ~/homework2/test
File Edit View Terminal Help

kbkim@ubuntu:~/homework2/test$ ls -i a*
931495 a
kbkim@ubuntu:~/homework2/test$ ln -s a a_sym
kbkim@ubuntu:~/homework2/test$ ln a a_hard
kbkim@ubuntu:~/homework2/test$ cp a a_cp
kbkim@ubuntu:~/homework2/test$ ls -i a*
931495 a 931651 a_cp 931495 a_hard 931649 a_sym
kbkim@ubuntu:~/homework2/test$ ls -l a_sym
lrwxrwxrwx 1 kbkim kbkim 1 2012-04-12 04:57 a_sym -> a
kbkim@ubuntu:~/homework2/test$ echo 'what!' > a_hard
kbkim@ubuntu:~/homework2/test$ cat a
what!
kbkim@ubuntu:~/homework2/test$ cat a_cp
This is a test
kbkim@ubuntu:~/homework2/test$ rm a
kbkim@ubuntu:~/homework2/test$ cat a_sym
cat: a_sym: No such file or directory
kbkim@ubuntu:~/homework2/test$ ls -i a*
931651 a_cp 931495 a_hard 931649 a_sym
kbkim@ubuntu:~/homework2/test$ ls -l a_sym
lrwxrwxrwx 1 kbkim kbkim 1 2012-04-12 04:57 a_sym -> a
kbkim@ubuntu:~/homework2/test$
```

# Problems of Filesystem

- Over time, an active filesystem can develop problems
  - It can fill up, causing individual programs or even the entire system to fail
  - It can become corrupted, perhaps due to a power failure or a system crash
  - It can run out of space for inodes, so no new files or directories can be created
- Monitoring and checking filesystems regularly can help prevent and correct problems

# Monitoring space : df

- “df” get a list of **free space** on all mounted filesystems
- “-h” option → human-readable option
- “Use%” column shows what percentage of the filesystem is in use
- Directory argument → make it show space on the filesystems those directories are mounted on
- “-i” option → inode usage



# Example of DF

```
kbkim@ubuntu: ~/testmount
File Edit View Terminal Help

kbkim@ubuntu:~/testmount$ df
Filesystem            1K-blocks      Used Available Use% Mounted on
/dev/sda1             19737268    2670648   16064024   15% /
none                  508728        276     508452    1% /dev
none                  512980        264     512716    1% /dev/shm
none                  512980        100     512880    1% /var/run
none                  512980         0     512980    0% /var/lock
none                  512980         0     512980    0% /lib/init/rw
:host:/               467664892   55650148  412014744   12% /mnt/hgfs

kbkim@ubuntu:~/testmount$ df -h
Filesystem            Size  Used Avail Use% Mounted on
/dev/sda1             19G   2.6G   16G   15% /
none                  497M   276K   497M    1% /dev
none                  501M   264K   501M    1% /dev/shm
none                  501M   100K   501M    1% /var/run
none                  501M    4.0K   501M    1% /var/lock
none                  501M     0     501M    0% /lib/init/rw
:host:/               446G   54G   393G   12% /mnt/hgfs

kbkim@ubuntu:~/testmount$ df .
Filesystem            1K-blocks      Used Available Use% Mounted on
/dev/sda1             19737268    2670648   16064024   15% /

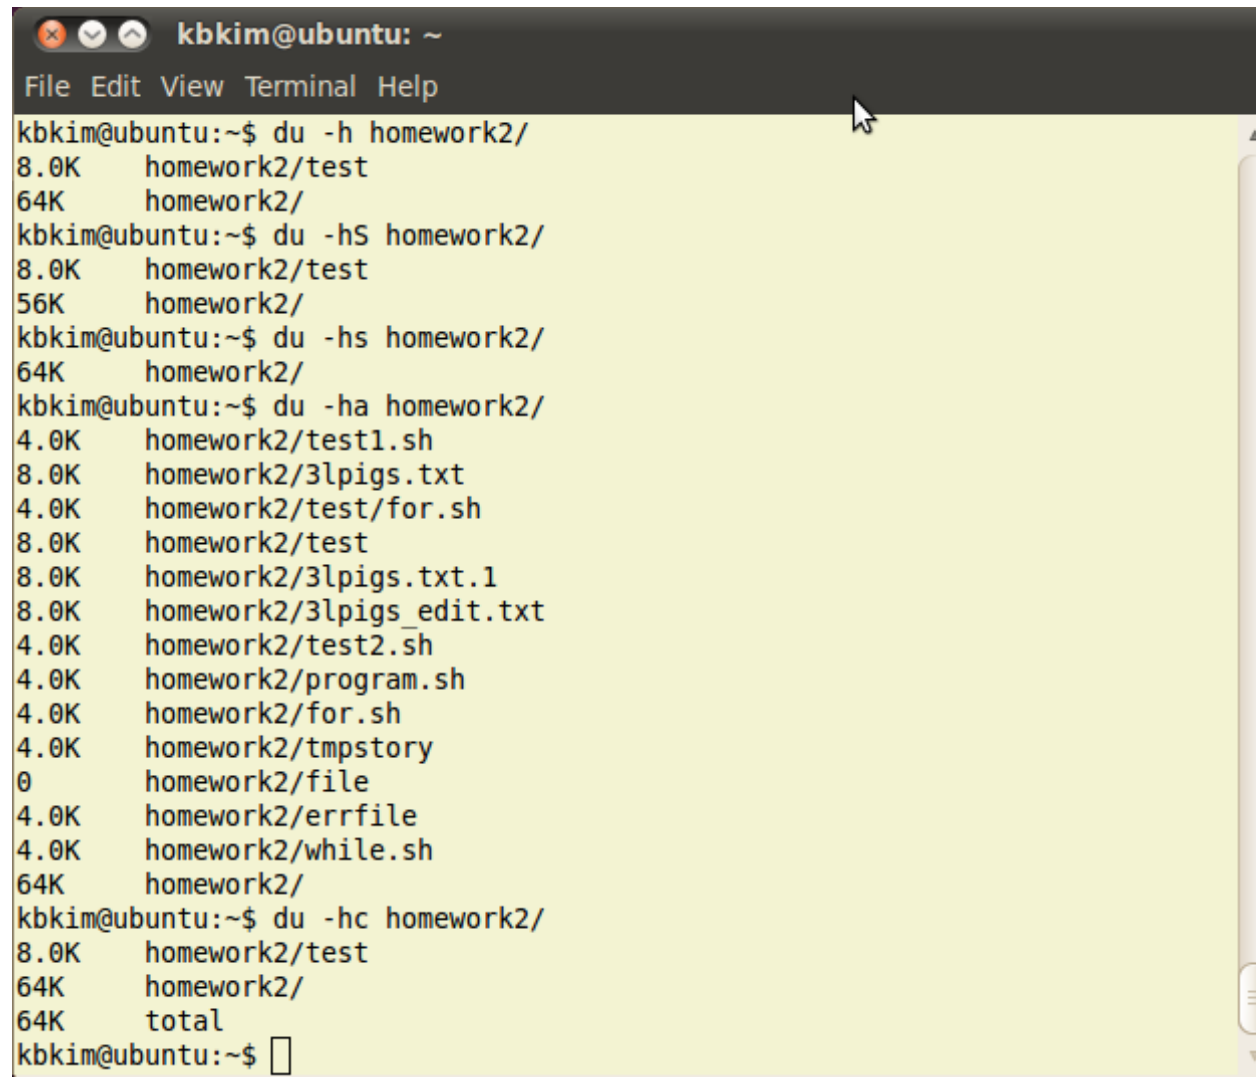
kbkim@ubuntu:~/testmount$ df -i
Filesystem            Inodes   IUsed   IFree IUse% Mounted on
/dev/sda1            1253376  143599 1109777   12% /
none                 127182     729  126453    1% /dev
none                 128245       7  128238    1% /dev/shm
none                 128245     52  128193    1% /var/run
none                 128245      3  128242    1% /var/lock
none                 128245      1  128244    1% /lib/init/rw
:host:/              0         0       0    - /mnt/hgfs

kbkim@ubuntu:~/testmount$
```

# Monitoring disk usage : du

- “du” shows a summary of the disk usage in a directory tree
- Options
  - “-a” → show all files
  - “-c” → cumulated total for all directories
  - “-h” → human-readable
  - “-S” → exclude subdirectories
  - “-s” → summary for each directory

# Example of du

A terminal window titled 'kbbkim@ubuntu: ~' with a menu bar containing 'File', 'Edit', 'View', 'Terminal', and 'Help'. The terminal shows the execution of the 'du' command with various options: '-h', '-hS', '-hs', and '-ha'. The output lists file sizes in human-readable format (K, M, G) for a directory named 'homework2' and its subdirectories and files. The window has a yellow background and a scrollbar on the right side.

```
kbbkim@ubuntu: ~  
File Edit View Terminal Help  
kbbkim@ubuntu:~$ du -h homework2/  
8.0K    homework2/test  
64K     homework2/  
kbbkim@ubuntu:~$ du -hS homework2/  
8.0K    homework2/test  
56K     homework2/  
kbbkim@ubuntu:~$ du -hs homework2/  
64K     homework2/  
kbbkim@ubuntu:~$ du -ha homework2/  
4.0K    homework2/test1.sh  
8.0K    homework2/3lpigs.txt  
4.0K    homework2/test/for.sh  
8.0K    homework2/test  
8.0K    homework2/3lpigs.txt.1  
8.0K    homework2/3lpigs_edit.txt  
4.0K    homework2/test2.sh  
4.0K    homework2/program.sh  
4.0K    homework2/for.sh  
4.0K    homework2/tmpstory  
0       homework2/file  
4.0K    homework2/errfile  
4.0K    homework2/while.sh  
64K     homework2/  
kbbkim@ubuntu:~$ du -hc homework2/  
8.0K    homework2/test  
64K     homework2/  
64K     total  
kbbkim@ubuntu:~$
```

# Finding and repairing filesystem : fsck

- Checks the integrity of a filesystem
  - Can make repairs if necessary
- Two main parts
  - Driver program, fsck
    - Handles any filesystem type
  - Backend program
    - For each specific filesystem type
    - Ext2 → e2fsck
      - Invoked through fsck
- “-f” option → force checking the filesystem
- Usually it is a bad idea to run fsck on a mounted filesystem

# What are Quotas?

- Quotas are a way of limiting the amount of disk space that users may take up
- Some organizations absolutely need to ensure that
  - No user can prevent other users from using a reasonable amount of disk space
  - No user can impede the correct functioning of the system

# Hard and soft limits

- Quotas have **hard limits** and **soft limits**
- A user can exceed the soft limit without retribution
  - But only for a certain period of time
    - The grace period
    - Set per-filesystem
  - The user is also warned that the soft limit has been exceeded
- A hard limit may never be exceeded
- If a user tries to exceed a hard limit, the attempt fails

# Block and Inode Limits

- Quotas can be set for blocks
  - Limits the amount of data space that may be used
- Quotas can be set for inodes
  - Limits the number of files that may be created

# Per-user and Per-group Quotas

- Most quotas are set per-user
  - Each user has his or her own soft limit and hard limit
- Quotas can also be set per-group
  - A group can be given a soft limit and hard limit
- Group quotas apply to all users in a group
- If a group hard limit has been reached, no user in the group may use more space
  - Including users who have not yet reached their individual quota



# Options in /etc/fstab

- The options in /etc/fstab specify which filesystems should have quota enabled
  - Add the option *usrquota* to enable user quotas
  - Use *grpquota* to enable group quotas
  - Either or both can be used for each filesystem
- The filesystems with quota enabled should have files called **aquota.user** and **aquota.group** in their root directories
- Mount option
  - “usrquota” and “grpquota”

# quotacheck command

- Check the users and groups related to quota regulation
- Option
  - “-c” → creating aquota db
  - “-u” → for user
  - “-g” → for group
  - “-f” → force checking
- “aquota.group” → db for group quota
- “aquota.user” → db for user quota

# setquota command

- Command line program to alter quota limits for a user or group
- Specify the name of a user or group with `-u username` or `-g groupname`
- Specify the filesystem
- Specify the limit to set must be specified in the following order
  - Soft limit for blocks
  - Hard limit for blocks
  - Soft limit for inodes
  - Hard limit for inodes
- Setting any limit to 0 will remove that limit
- `-t` option : set grace period of block and inode softlimit in seconds
- Examples
  - `setquota -u kbkim 10 20 10 20 /dev/hda1`
  - `setquota -t 3600 3600 /dev/hda1`

# repquota/quotanon/quotaooff commands

- “quotaon”/”quotaoff”
  - Turns on/off quota support
  - Can only be done by root
- “Repquota”
  - Prints information about the quota limits assigned to each user
  - Root privilege
  - “-a” → all filesystems
  - “-g” → group quotas
  - “-v” → more complete information
- “quota”
  - For each users to check their quota

# Example of quota

```
kbkim@ubuntu: ~/testmount
File Edit View Terminal Help

root@ubuntu:/home/kbkim/testmount# mount -o loop,usrquota,grpquota disk.img disk
root@ubuntu:/home/kbkim/testmount# ls disk
lost+found
root@ubuntu:/home/kbkim/testmount# repquota disk
repquota: Quota file not found or has wrong format.
repquota: Not all specified mountpoints are using quota.
root@ubuntu:/home/kbkim/testmount# quotacheck -cug disk
root@ubuntu:/home/kbkim/testmount# ls disk
aquota.group aquota.user lost+found
root@ubuntu:/home/kbkim/testmount# repquota disk
*** Report for user quotas on device /dev/loop0
Block grace time: 7days; Inode grace time: 7days

```

		Block limits				File limits			
User	used	soft	hard	grace	used	soft	hard	grace	
root	-- 18	0	0		2	0	0		
kbkim	-- 1	0	0		1	0	0		

```

root@ubuntu:/home/kbkim/testmount# quota kbkim
Disk quotas for user kbkim (uid 1000): none
root@ubuntu:/home/kbkim/testmount# su kbkim
kbkim@ubuntu:~/testmount$ quota kbkim
quota: Cannot open quotafile /home/kbkim/testmount/disk/aquota.user: Permission denied
quota: Quota file not found or has wrong format.
kbkim@ubuntu:~/testmount$
```

# Example of setquota

```
kbkim@ubuntu: ~/testmount
File Edit View Terminal Help

root@ubuntu:/home/kbkim/testmount# quotaon disk
root@ubuntu:/home/kbkim/testmount# su kbkim
kbkim@ubuntu:~/testmount$ quota kbkim
Disk quotas for user kbkim (uid 1000): none
kbkim@ubuntu:~/testmount$ exit
exit
root@ubuntu:/home/kbkim/testmount# setquota -u kbkim 3 5 3 5 disk
root@ubuntu:/home/kbkim/testmount# setquota -t 30 30 disk
root@ubuntu:/home/kbkim/testmount# repquota disk
*** Report for user quotas on device /dev/loop0
Block grace time: 00:01; Inode grace time: 00:01
```

		Block limits				File limits			
User	used	soft	hard	grace	used	soft	hard	grace	
root	--	18	0	0	2	0	0		
kbkim	--	1	3	5	1	3	5		

```

root@ubuntu:/home/kbkim/testmount# quota kbkim
Disk quotas for user kbkim (uid 1000):
    Filesystem blocks quota limit grace files quota limit grace
    /dev/loop0    1      3      5      1      3      5
root@ubuntu:/home/kbkim/testmount# su kbkim
kbkim@ubuntu:~/testmount$ quota
Disk quotas for user kbkim (uid 1000):
    Filesystem blocks quota limit grace files quota limit grace
    /dev/loop0    1      3      5      1      3      5
kbkim@ubuntu:~/testmount$
```

# Example of quota

```
kbkim@ubuntu: ~/testmount
File Edit View Terminal Help

kbkim@ubuntu:~/testmount$ touch disk/1
kbkim@ubuntu:~/testmount$ quota
Disk quotas for user kbkim (uid 1000):
  Filesystem blocks quota limit grace files quota limit grace
    /dev/loop0      1    3    5
kbkim@ubuntu:~/testmount$ touch disk/2
kbkim@ubuntu:~/testmount$ quota
Disk quotas for user kbkim (uid 1000):
  Filesystem blocks quota limit grace files quota limit grace
    /dev/loop0      1    3    5
kbkim@ubuntu:~/testmount$ touch disk/3
kbkim@ubuntu:~/testmount$ quota
Disk quotas for user kbkim (uid 1000):
  Filesystem blocks quota limit grace files quota limit grace
    /dev/loop0      1    3    5
kbkim@ubuntu:~/testmount$ touch disk/4
kbkim@ubuntu:~/testmount$ quota
Disk quotas for user kbkim (uid 1000):
  Filesystem blocks quota limit grace files quota limit grace
    /dev/loop0      1    3    5
kbkim@ubuntu:~/testmount$ touch disk/5
touch: cannot touch `disk/5': Disk quota exceeded
kbkim@ubuntu:~/testmount$ quota
Disk quotas for user kbkim (uid 1000):
  Filesystem blocks quota limit grace files quota limit grace
    /dev/loop0      1    3    5
kbkim@ubuntu:~/testmount$ rm disk/4
kbkim@ubuntu:~/testmount$ quota
Disk quotas for user kbkim (uid 1000):
  Filesystem blocks quota limit grace files quota limit grace
    /dev/loop0      1    3    5
kbkim@ubuntu:~/testmount$ touch disk/4
touch: cannot touch `disk/4': Disk quota exceeded
kbkim@ubuntu:~/testmount$
```